

Fast matrix multiplication using CP decomposition

Computational laboratory project report

Alberto Defendi

Abstract

Bla bla bla

Contents

1	Notation and tensor preliminaries	1
2	Low rank tensor decomposition	1
3	Fast matrix multiplication	1
4	Implementation and practical consideration	3
5	Algorithm schema	3
6	Dynamic Peeling	3
7	Parallel algorithms for shared memory	4
7.1	Implementation	4
7.2	Benchmarking	4

1 Notation and tensor preliminaries

A tensor is a multidimensional array, and in this work we deal only with 3-dimensional (or 3-way) tensors $\mathcal{X} \in \mathbb{R}^{m \times n \times p}$.

Short recap
of tensor
stuff

2 Low rank tensor decomposition

3 Fast matrix multiplication

Matrix multiplication is a bilinear operation. We indicate with $\langle m, n, p \rangle$ the bilinear form that represents the multiplication between a matrix $\mathbf{A} \in \mathbb{R}^{m \times n}$ with $\mathbf{B} \in \mathbb{R}^{n \times p}$. Fixing the natural ordering on the

entries of $\mathbf{A}$ and $\mathbf{B}$ and the inverse ordering on $\mathbf{C} = \mathbf{AB}$, we get a representation of $\mathfrak{X} \in \mathbb{R}^{mn \times np \times mp}$ of the matrix multiplication $\langle m, n, p \rangle$ such that

$$(1) \quad \text{vec}(\mathbf{C}^\top) = \mathfrak{X} \times_1 \text{vec}(\mathbf{A})^\top \times_2 \text{vec}(\mathbf{B})^\top.$$

Note that we fix the linear order on A and B , while we fix the reverse ordering on C by applying the transpose operation.

If we are given a CP decomposition of rank r of the tensor $\mathfrak{X} = \llbracket U, V, W \rrbracket$, we can rewrite Equation 1 into

$$(2) \quad \text{vec}(\mathbf{C}^\top) = \sum_{l=1}^r (\underbrace{\mathbf{u}_l^\top \text{vec}(\mathbf{A})}_{s_l}) \cdot (\underbrace{\mathbf{v}_l^\top \text{vec}(\mathbf{B})}_{t_l}) \mathbf{w}_l = \sum_{l=1}^r (\underbrace{s_l \cdot t_l}_{m_l}) w_l,$$

where $\mathbf{u}_l, \mathbf{v}_l, \mathbf{w}_l$ denote the columns of $\mathbf{U}, \mathbf{V}$ and $\mathbf{W}$. This reduces the number of active multiplications to the rank r of $\mathfrak{X}$.

When $m, n, p > 2$ we can always partition them in blocks as:

$$(3) \quad \begin{matrix} \lfloor m/2 \rfloor \{ & \overbrace{\begin{bmatrix} A_{11} & A_{12} \\ A_{21} & A_{22} \end{bmatrix}}^{\lfloor n/2 \rfloor \times \lfloor n/2 \rfloor} & \} \\ \lceil m/2 \rceil \{ & \overbrace{\begin{bmatrix} B_{11} & B_{12} \\ B_{21} & B_{22} \end{bmatrix}}^{\lfloor p/2 \rfloor \times \lfloor p/2 \rfloor} & \} \end{matrix},$$

where $A_{11} \in \mathbb{R}^{\lfloor \frac{m}{2} \rfloor \times \lfloor \frac{n}{2} \rfloor}$, $B_{11} \in \mathbb{R}^{\lfloor \frac{n}{2} \rfloor \times \lfloor \frac{p}{2} \rfloor}$, $A_{22} \in \mathbb{R}^{\lceil \frac{m}{2} \rceil \times \lceil \frac{n}{2} \rceil}$, and $B_{22} \in \mathbb{R}^{\lceil \frac{n}{2} \rceil \times \lceil \frac{p}{2} \rceil}$.

In particular, we can view the product AB as the product between 2×2 block matrices

$$(4) \quad \begin{bmatrix} C_{11} & C_{12} \\ C_{21} & C_{22} \end{bmatrix} = \begin{bmatrix} A_{11} & A_{12} \\ A_{21} & A_{22} \end{bmatrix} \cdot \begin{bmatrix} B_{11} & B_{12} \\ B_{21} & B_{22} \end{bmatrix} = \begin{bmatrix} A_{11}B_{11} + A_{12}B_{21} & A_{11}B_{12} + A_{12}B_{22} \\ A_{21}B_{11} + A_{22}B_{21} & A_{21}B_{12} + A_{22}B_{22} \end{bmatrix}.$$

The naive algorithm proceeds by combining a set of eight matrix multiplications with four additions:

$$\begin{aligned} M_1 &= A_{11} \cdot B_{11} & M_2 &= A_{12} \cdot B_{21} & M_3 &= A_{11} \cdot B_{12} \\ M_4 &= A_{12} \cdot B_{22} & M_5 &= A_{21} \cdot B_{11} & M_6 &= A_{22} \cdot B_{21} \\ M_7 &= A_{21} \cdot B_{12} & M_8 &= A_{22} \cdot B_{22} \end{aligned}$$

$$\begin{aligned} C_{11} &= M_1 + M_2 & C_{12} &= M_3 + M_4 \\ C_{21} &= M_5 + M_6 & C_{22} &= M_7 + M_8 \end{aligned}$$

For example, when $m = n = p = 2$, we get the tensor with frontal slices

$$\mathbf{X}_1 = \begin{bmatrix} 1 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \\ 0 & 0 & 0 & 0 \end{bmatrix}, \quad \mathbf{X}_2 = \begin{bmatrix} 0 & 0 & 1 & 0 \\ 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 1 \\ 0 & 0 & 0 & 0 \end{bmatrix}, \quad \mathbf{X}_3 = \begin{bmatrix} 0 & 0 & 0 & 0 \\ 1 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \end{bmatrix}, \quad \mathbf{X}_4 = \begin{bmatrix} 0 & 0 & 0 & 0 \\ 0 & 0 & 1 & 0 \\ 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 1 \end{bmatrix}.$$

why

Tensor
definition

This yields, for example,

$$X_3 \times_1 \text{vec}(A)^\top \times_2 \text{vec}(B)^\top = \begin{bmatrix} a_{11} & a_{21} & a_{12} & a_{22} \end{bmatrix} \begin{bmatrix} 0 & 0 & 0 & 0 \\ 1 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \end{bmatrix} \begin{bmatrix} b_{11} \\ b_{21} \\ b_{12} \\ b_{22} \end{bmatrix} = a_{21}b_{11} + a_{22}b_{21} = c_{21}.$$

Note that the formula in Equation 2 in the context of block matrix multiplication in Equation 4 yields to

$$\text{vec}(A) = \begin{bmatrix} A_{11} \\ A_{21} \\ A_{12} \\ A_{22} \end{bmatrix}, \quad \text{vec}(B) = \begin{bmatrix} B_{11} \\ B_{21} \\ B_{12} \\ B_{22} \end{bmatrix}, \quad \text{vec}(C^\top) = \begin{bmatrix} C_{11} \\ C_{12} \\ C_{21} \\ C_{22} \end{bmatrix}.$$

Also the factors s_l and t_l in Equation 2 become

$$s_l = u_l^\top \text{vec}(A) = \sum_{i=1}^4 U_{li} \text{vec}(A)_i, \quad t_l = V_l^\top \text{vec}(B) = \sum_{i=1}^4 V_{li} \text{vec}(B)_i$$

4 Implementation and practical consideration

Following the work [1] we implement a fast multiplication algorithm based on CP decomposition given the U , V and W matrices representing the algorithm.

Fast matrix multiplication algorithm such as Strassen face two significant problems in practical implementation, which are recursion cutoff and input matrices sizes not matching with base case. The former can solved by ...and the latter by dynamic peeling 6.

5 Algorithm schema

6 Dynamic Peeling

In Strassen algorithm we want to multiply $\langle m, n, p \rangle$ matrices by splitting in four blocks as in Equation 3. If m , n and p are even, then it is not possible to split blocks in even sizes. The easiest fix for this is padding the matrix with zeros until an even size is reached, but this adds a significant memory waste. A more efficient approach, called *dynamic peeling*, consists in stripping the extra row off and/or column and adding their contribution to the final results. In detail,

For the matrix multiplication $C = AB$, where A is an $m \times n$ matrix and B is an $n \times p$ matrix. The matrix A is split into an $(m-1) \times (n-1)$ core matrix A_{11} , alongside its peeled boundary vectors a_{12} , a_{21} , and the scalar corner element a_{22} . Similarly, the matrix B is split into an $(n-1) \times (p-1)$ core matrix B_{11} , with its corresponding peeled components b_{12} , b_{21} , and b_{22} :

$$A = \left[\begin{array}{c|c} \overbrace{A_{11}}^{n-1} & \overbrace{a_{12}}^1 \\ \hline a_{21} & a_{22} \end{array} \right] \left. \vphantom{\begin{array}{c|c} \overbrace{A_{11}}^{n-1} & \overbrace{a_{12}}^1 \\ \hline a_{21} & a_{22} \end{array}} \right\}^{m-1} \vphantom{\begin{array}{c|c} \overbrace{A_{11}}^{n-1} & \overbrace{a_{12}}^1 \\ \hline a_{21} & a_{22} \end{array}} \Bigg\}^1, \quad B = \left[\begin{array}{c|c} \overbrace{B_{11}}^{p-1} & \overbrace{b_{12}}^1 \\ \hline b_{21} & b_{22} \end{array} \right] \left. \vphantom{\begin{array}{c|c} \overbrace{B_{11}}^{p-1} & \overbrace{b_{12}}^1 \\ \hline b_{21} & b_{22} \end{array}} \right\}^{n-1} \vphantom{\begin{array}{c|c} \overbrace{B_{11}}^{p-1} & \overbrace{b_{12}}^1 \\ \hline b_{21} & b_{22} \end{array}} \Bigg\}^1$$

The final block matrix product is computed directly through a combination of the core Strassen multiplication ($A_{11}B_{11}$) and low-rank vector/scalar fixup modifications:

$$\left[\begin{array}{c|c} \overbrace{C_{11}}^{p-1} & \overbrace{c_{12}}^1 \\ \hline c_{21} & c_{22} \end{array} \right] \left. \vphantom{\begin{array}{c|c} \overbrace{C_{11}}^{p-1} & \overbrace{c_{12}}^1 \\ \hline c_{21} & c_{22} \end{array}} \right\}^{m-1} \vphantom{\begin{array}{c|c} \overbrace{C_{11}}^{p-1} & \overbrace{c_{12}}^1 \\ \hline c_{21} & c_{22} \end{array}} \Bigg\}^1 = \left[\begin{array}{c|c} A_{11}B_{11} + a_{12}b_{21} & A_{11}b_{12} + a_{12}b_{22} \\ \hline a_{21}B_{11} + a_{22}b_{21} & a_{21}b_{12} + a_{22}b_{22} \end{array} \right]$$

Here, only the sub-product $A_{11}B_{11}$ is executed recursively using Strassen's matrix multiplication, while all other terms are computed via standard matrix multiplication.

7 Parallel algorithms for shared memory

7.1 Implementation

We use Rust's Rayon to handle parallelization, and we can summarize the parallel scheduling as following:

- **BFS:** Each recursive matrix multiplication routine and the associated matrix additions are launched by Rayon as an iterable, and then are collected in an array of matrices $[M_1, \dots, M_r]$.
- **DFS:** Each Faer matrix multiplication call uses all thread, and matrix multiplications are always fully parallelized.
- **Hybrid:** TODO.

implement

7.2 Benchmarking

All benchmarks are run with Rust's Criterion library which performs warm-up before doing measurements. In order to compare the performance of matrix multiplication algorithms with different computational costs, we use the *effective* GFLOPS metric for $\langle m, n, r \rangle$ matrix multiplication:

$$(5) \quad \text{effective GFLOPS} = \frac{2mnp - mp}{\text{time in seconds} \cdot 10^9},$$

where the numerator can be rewritten as $2mnp - mp = mnp + m(n-1)p$, which is the number of multiplications and additions performed for classical matrix multiplication. This allows us to compare all of the algorithms on an inverse-time scale, normalized by problem size [1]. For fast algorithms, considering both multiplications and additions in the metric is more accurate than taking multiplications only [2].

References

- [1] Austin R. Benson and Grey Ballard. A framework for practical parallel fast matrix multiplication. In *Proceedings of the 20th ACM SIGPLAN Symposium on Principles and Practice of Parallel Programming*, PPOPP 2015, page 42–53, New York, NY, USA, 2015. Association for Computing Machinery.
- [2] Benjamin Lipshitz, Grey Ballard, James Demmel, and Oded Schwartz. Communication-avoiding parallel strassen: Implementation and performance. In *SC '12: Proceedings of the International Conference on High Performance Computing, Networking, Storage and Analysis*, pages 1–11, 2012.
- [3] S. Huss-Lederman, E.M. Jacobson, J.R. Johnson, A. Tsao, and T. Turnbull. Implementation of strassen’s algorithm for matrix multiplication. In *Supercomputing '96: Proceedings of the 1996 ACM/IEEE Conference on Supercomputing*, pages 32–32, 1996.